

# Order Delivery Time Prediction

## Objectives

The objective of this assignment is to build a regression model that predicts the delivery time for orders placed through Porter. The model will use various features such as the items ordered, the restaurant location, the order protocol, and the availability of delivery partners.

The key goals are:

- Predict the delivery time for an order based on multiple input features
- Improve delivery time predictions to optimiae operational efficiency
- Understand the key factors influencing delivery time to enhance the model's accuracy

## Data Pipeline

The data pipeline for this assignment will involve the following steps:

1. **Data Loading**
2. **Data Preprocessing and Feature Engineering**
3. **Exploratory Data Analysis**
4. **Model Building**
5. **Model Inference**

## Data Understanding

The dataset contains information on orders placed through Porter, with the following columns:

Field	Description
market_id	Integer ID representing the market where the restaurant is located.
created_at	Timestamp when the order was placed.
actual_delivery_time	Timestamp when the order was delivered.
store_primary_category	Category of the restaurant (e.g., fast food, dine-in).
order_protocol	Integer representing how the order was placed (e.g., via Porter, call to restaurant, etc.).
total_items	Total number of items in the order.
subtotal	Final price of the order.
num_distinct_items	Number of distinct items in the order.
min_item_price	Price of the cheapest item in the order.
max_item_price	Price of the most expensive item in the order.

Field	Description
total_onshift_dashers	Number of delivery partners on duty when the order was placed.
total_busy_dashers	Number of delivery partners already occupied with other orders.
total_outstanding_orders	Number of orders pending fulfillment at the time of the order.
distance	Total distance from the restaurant to the customer.

## Importing Necessary Libraries

```
# Import essential libraries for data manipulation and analysis

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

pd.set_option('display.max_columns', None, 'display.max_rows', None)
# Show all columns in DataFrame
# Set the style for seaborn plots
sns.set(style='whitegrid')
```

## 1. Loading the data

Load 'porter\_data\_1.csv' as a DataFrame

```
# Importing the file porter_data_1.csv
df = pd.read_csv('porter_data_1.csv')
```

## 2. Data Preprocessing and Feature Engineering [15 marks]

### 2.1 Fixing the Datatypes [5 marks]

The current timestamps are in object format and need conversion to datetime format for easier handling and intended functionality

#### 2.1.1 [2 marks]

Convert date and time fields to appropriate data type

```
df.head()
```

	market_id	created_at	actual_delivery_time	\
0	1.0	2015-02-06 22:24:17	2015-02-06 23:11:17	
1	2.0	2015-02-10 21:49:25	2015-02-10 22:33:25	

2	2.0	2015-02-16 00:11:35	2015-02-16 01:06:35
3	1.0	2015-02-12 03:36:46	2015-02-12 04:35:46
4	1.0	2015-01-27 02:12:36	2015-01-27 02:58:36

	store_primary_category	order_protocol	total_items	subtotal	\
0	4	1.0	4	3441	
1	46	2.0	1	1900	
2	36	3.0	4	4771	
3	38	1.0	1	1525	
4	38	1.0	2	3620	

	num_distinct_items	min_item_price	max_item_price
total_onshift_dashers	\		
0	4	557	1239
33.0			
1	1	1400	1400
1.0			
2	3	820	1604
8.0			
3	1	1525	1525
5.0			
4	2	1425	2195
5.0			

	total_busy_dashers	total_outstanding_orders	distance
0	14.0	21.0	34.44
1	2.0	2.0	27.60
2	6.0	18.0	11.56
3	6.0	8.0	31.80
4	5.0	7.0	8.20

*# Convert 'created\_at' and 'actual\_delivery\_time' columns to datetime format*

```
df['created_at'] = pd.to_datetime(df['created_at'])
df['actual_delivery_time'] =
pd.to_datetime(df['actual_delivery_time'])
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 175777 entries, 0 to 175776
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	market_id	175777 non-null	float64
1	created_at	175777 non-null	datetime64[ns]
2	actual_delivery_time	175777 non-null	datetime64[ns]
3	store_primary_category	175777 non-null	int64
4	order_protocol	175777 non-null	float64

```

5  total_items      175777 non-null  int64
6  subtotal        175777 non-null  int64
7  num_distinct_items 175777 non-null  int64
8  min_item_price   175777 non-null  int64
9  max_item_price   175777 non-null  int64
10 total_onshift_dashers 175777 non-null float64
11 total_busy_dashers 175777 non-null float64
12 total_outstanding_orders 175777 non-null float64
13 distance         175777 non-null float64
dtypes: datetime64[ns](2), float64(6), int64(6)
memory usage: 18.8 MB

```

total\_busy\_dashers, total\_busy\_dashers, total\_outstanding\_orders should be in int64 not float64, can't be in decimal

```

df[['total_onshift_dashers', 'total_busy_dashers',
'total_outstanding_orders', 'market_id', 'order_protocol']] =
df[['total_onshift_dashers', 'total_busy_dashers',
'total_outstanding_orders', 'market_id', 'order_protocol']].astype('int64')

```

```
df.head()
```

	market_id	created_at	actual_delivery_time
store_primary_category \			
0	1	2015-02-06 22:24:17	2015-02-06 23:11:17
4			
1	2	2015-02-10 21:49:25	2015-02-10 22:33:25
46			
2	2	2015-02-16 00:11:35	2015-02-16 01:06:35
36			
3	1	2015-02-12 03:36:46	2015-02-12 04:35:46
38			
4	1	2015-01-27 02:12:36	2015-01-27 02:58:36
38			

	order_protocol	total_items	subtotal	num_distinct_items
min_item_price \				
0	1	4	3441	4
557				
1	2	1	1900	1
1400				
2	3	4	4771	3
820				
3	1	1	1525	1
1525				
4	1	2	3620	2
1425				

	max_item_price	total_onshift_dashers	total_busy_dashers	\
--	----------------	-----------------------	--------------------	---

0	1239	33	14
1	1400	1	2
2	1604	8	6
3	1525	5	6
4	2195	5	5

	total_outstanding_orders	distance
0	21	34.44
1	2	27.60
2	18	11.56
3	8	31.80
4	7	8.20

### 2.1.2 [3 marks]

Convert categorical fields to appropriate data type

```
df.market_id.value_counts().sort_index()

market_id
1      37115
2      53469
3      21075
4      46222
5      17258
6        638
Name: count, dtype: int64

# Convert categorical features to category type

df[['market_id', 'store_primary_category', 'order_protocol']] =
df[['market_id', 'store_primary_category', 'order_protocol']].astype('category')
```

## 2.2 Feature Engineering [5 marks]

Calculate the time taken to execute the delivery as well as extract the hour and day at which the order was placed

### 2.2.1 [2 marks]

Calculate the time taken using the features `actual_delivery_time` and `created_at`

```
# Calculate time taken in minutes
df['mins_taken'] = (df['actual_delivery_time'] -
df['created_at']).dt.total_seconds() / 60

df.head()

market_id      created_at  actual_delivery_time
store_primary_category \
```

```

0      1 2015-02-06 22:24:17 2015-02-06 23:11:17
4
1      2 2015-02-10 21:49:25 2015-02-10 22:33:25
46
2      2 2015-02-16 00:11:35 2015-02-16 01:06:35
36
3      1 2015-02-12 03:36:46 2015-02-12 04:35:46
38
4      1 2015-01-27 02:12:36 2015-01-27 02:58:36
38

```

```

order_protocol  total_items  subtotal  num_distinct_items
min_item_price \
0              1           4       3441           4
557
1              2           1       1900           1
1400
2              3           4       4771           3
820
3              1           1       1525           1
1525
4              1           2       3620           2
1425

```

```

max_item_price  total_onshift_dashers  total_busy_dashers \
0             1239                   33                   14
1             1400                    1                    2
2             1604                    8                    6
3             1525                    5                    6
4             2195                    5                    5

```

```

total_outstanding_orders  distance  mints_taken
0                        21      34.44      47.0
1                         2      27.60      44.0
2                        18      11.56      55.0
3                         8      31.80      59.0
4                         7       8.20      46.0

```

### 2.2.2 [3 marks]

Extract the hour at which the order was placed and which day of the week it was. Drop the unnecessary columns.

```

# Extract the hour and day of week from the 'created_at' timestamp

df['created_hour'] = df['created_at'].dt.hour
df['is_weekend'] = df['created_at'].dt.dayofweek.apply(lambda x: 0 if
x < 5 else 1)
df['day_of_week'] = df['created_at'].dt.dayofweek
# Create a categorical feature 'isWeekend'

```

```
df['is_weekend'] = df['is_weekend'].astype('category')
```

```
df.head()
```

	market_id	created_at	actual_delivery_time
0	1	2015-02-06 22:24:17	2015-02-06 23:11:17
4			
1	2	2015-02-10 21:49:25	2015-02-10 22:33:25
46			
2	2	2015-02-16 00:11:35	2015-02-16 01:06:35
36			
3	1	2015-02-12 03:36:46	2015-02-12 04:35:46
38			
4	1	2015-01-27 02:12:36	2015-01-27 02:58:36
38			

	order_protocol	total_items	subtotal	num_distinct_items
0	1	4	3441	4
557				
1	2	1	1900	1
1400				
2	3	4	4771	3
820				
3	1	1	1525	1
1525				
4	1	2	3620	2
1425				

	max_item_price	total_onshift_dashers	total_busy_dashers
0	1239	33	14
1	1400	1	2
2	1604	8	6
3	1525	5	6
4	2195	5	5

	total_outstanding_orders	distance	mints_taken	created_hour
0	21	34.44	47.0	22
0				
1	2	27.60	44.0	21
0				
2	18	11.56	55.0	0
0				
3	8	31.80	59.0	3
0				
4	7	8.20	46.0	2
0				

	day_of_week
0	4
1	1
2	0
3	3
4	1

There's not much need to drop any column, maybe `actual_delivery_time` bcoz the prediction done on delivery time, best to do that duration from order `created_time`

```
# Drop unnecessary columns
df = df.drop(columns=['actual_delivery_time', 'created_at'])

print(df['store_primary_category'].value_counts(), df['order_protocol'].value_counts(), df['market_id'].value_counts())
```

store_primary_category	
4	18183
55	15745
46	15586
13	9915
58	8995
20	8563
39	8232
24	8085
38	6733
28	6495
36	6378
68	6235
72	5570
45	5138
10	4841
50	3714
57	3468
34	2951
7	2672
59	2590
6	2219
66	2103
15	2027
2	1745
40	1693
18	1550
61	1525
47	1457
35	1451
65	1032
25	992



71	755
14	639
52	625
53	574
30	557
42	490
12	428
23	330
16	310
9	293
49	287
29	259
17	243
70	232
54	230
69	220
31	182
51	139
11	125
67	125
26	110
0	103
44	92
5	70
32	66
62	57
37	55
41	51
63	41
60	30
64	25
19	24
33	24
48	24
22	23
27	22
56	11
1	10
43	9
8	2
3	1
21	1

Name: count, dtype: int64 order\_protocol

1	48404
3	47125
5	41415
2	20890
4	17246
6	678

```

7      19
Name: count, dtype: int64 market_id
2      53469
4      46222
1      37115
3      21075
5      17258
6       638
Name: count, dtype: int64

```

```

df = pd.get_dummies(columns=['order_protocol', 'market_id'],
prefix=['order_protocol', 'market_id'], data=df, drop_first=True)
df.head()

```

	store_primary_category	total_items	subtotal	num_distinct_items	\
0	4	4	3441	4	
1	46	1	1900	1	
2	36	4	4771	3	
3	38	1	1525	1	
4	38	2	3620	2	

	min_item_price	max_item_price	total_onshift_dashers
total_busy_dashers \			
0	557	1239	33
14			
1	1400	1400	1
2			
2	820	1604	8
6			
3	1525	1525	5
6			
4	1425	2195	5
5			

	total_outstanding_orders	distance	mints_taken	created_hour
is_weekend \				
0	21	34.44	47.0	22
0				
1	2	27.60	44.0	21
0				
2	18	11.56	55.0	0
0				
3	8	31.80	59.0	3
0				
4	7	8.20	46.0	2
0				

	day_of_week	order_protocol_2	order_protocol_3
order_protocol_4 \			
0	4	False	False
			False

1	1	True	False	False
2	0	False	True	False
3	3	False	False	False
4	1	False	False	False

	order_protocol_5	order_protocol_6	order_protocol_7	
market_id_2 \				
0	False	False	False	False
1	False	False	False	True
2	False	False	False	True
3	False	False	False	False
4	False	False	False	False

	market_id_3	market_id_4	market_id_5	market_id_6
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False

```

top_categories =
df['store_primary_category'].value_counts().nlargest(10).index
df['store_primary_category'] = df['store_primary_category'].apply(
    lambda x: x if x in top_categories else 'Other'
)
df = pd.get_dummies(df, columns=['store_primary_category'],
drop_first=True)

```

## 2.3 Creating training and validation sets [5 marks]

### 2.3.1 [2 marks]

Define target and input features

```
df.head()
```

	total_items	subtotal	num_distinct_items	min_item_price
max_item_price \				
0	4	3441	4	557
1239				
1	1	1900	1	1400

1400				
2	4	4771	3	820
1604				
3	1	1525	1	1525
1525				
4	2	3620	2	1425
2195				

	total_onshift_dashers	total_busy_dashers	total_outstanding_orders
\			
0	33	14	21
1	1	2	2
2	8	6	18
3	5	6	8
4	5	5	7

	distance	mints_taken	created_hour	is_weekend	day_of_week	\
0	34.44	47.0	22	0	4	
1	27.60	44.0	21	0	1	
2	11.56	55.0	0	0	0	
3	31.80	59.0	3	0	3	
4	8.20	46.0	2	0	1	

	order_protocol_2	order_protocol_3	order_protocol_4
order_protocol_5 \			
0	False	False	False
False			
1	True	False	False
False			
2	False	True	False
False			
3	False	False	False
False			
4	False	False	False
False			

	order_protocol_6	order_protocol_7	market_id_2	market_id_3
market_id_4 \				
0	False	False	False	False
False				
1	False	False	True	False
False				
2	False	False	True	False
False				
3	False	False	False	False

False				
4	False	False	False	False
False				

	market_id_5	market_id_6	store_primary_category_13	\
0	False	False	False	
1	False	False	False	
2	False	False	False	
3	False	False	False	
4	False	False	False	

	store_primary_category_20	store_primary_category_24	\
0		False	False
1		False	False
2		False	False
3		False	False
4		False	False

	store_primary_category_28	store_primary_category_38	\
0		False	False
1		False	False
2		False	False
3		False	True
4		False	True

	store_primary_category_39	store_primary_category_46	\
0		False	False
1		False	True
2		False	False
3		False	False
4		False	False

	store_primary_category_55	store_primary_category_58	\
0		False	False
1		False	False
2		False	False
3		False	False
4		False	False

	store_primary_category_Other
0	False
1	False
2	True
3	False
4	False

```
# Define target variable (y) and features (X)
y = df['mints_taken']
X = df.drop(columns=['mints_taken'])
```

```
X.head()
```

	total_items	subtotal	num_distinct_items	min_item_price
0	4	3441	4	557
1	1	1900	1	1400
2	4	4771	3	820
3	1	1525	1	1525
4	2	3620	2	1425

	total_onshift_dashers	total_busy_dashers	total_outstanding_orders
0	33	14	21
1	1	2	2
2	8	6	18
3	5	6	8
4	5	5	7

	distance	created_hour	is_weekend	day_of_week	order_protocol_2
0	34.44	22	0	4	False
1	27.60	21	0	1	True
2	11.56	0	0	0	False
3	31.80	3	0	3	False
4	8.20	2	0	1	False

	order_protocol_3	order_protocol_4	order_protocol_5
0	False	False	False
1	False	False	False
2	True	False	False
3	False	False	False
4	False	False	False

	order_protocol_7	market_id_2	market_id_3	market_id_4
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False

0	False	False	False	False
1	False	True	False	False
2	False	True	False	False
3	False	False	False	False
4	False	False	False	False

market_id_6 store_primary_category_13 store_primary_category_20 \				
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False

store_primary_category_24 store_primary_category_28 \			
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False

store_primary_category_38 store_primary_category_39 \			
0	False	False	False
1	False	False	False
2	False	False	False
3	True	False	False
4	True	False	False

store_primary_category_46 store_primary_category_55 \			
0	False	False	False
1	True	False	False
2	False	False	False
3	False	False	False
4	False	False	False

store_primary_category_58 store_primary_category_Other			
0	False	False	False
1	False	False	False
2	False	True	True

3	False	False
4	False	False

### 2.3.2 [3 marks]

Split the data into training and test sets

```
# Split data into training and testing sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y,
test_size=0.3, random_state=42)
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)

(123043, 33) (52734, 33) (123043,) (52734,)

X_train.head()
```

	total_items	subtotal	num_distinct_items	min_item_price	\
94465	3	4385	3	1095	
100712	4	2772	3	79	
153524	7	1663	5	139	
85660	3	1477	3	429	
100506	1	1134	1	725	

	max_item_price	total_onshift_dashers	total_busy_dashers	\
94465	1395	84	63	
100712	1349	22	16	
153524	579	2	1	
85660	579	41	42	
100506	725	60	61	

	total_outstanding_orders	distance	created_hour	is_weekend	\
94465	101	26.72	2	1	
100712	13	22.88	3	1	
153524	0	13.80	5	0	
85660	41	8.76	3	0	
100506	79	12.40	20	0	

	day_of_week	order_protocol_2	order_protocol_3
order_protocol_4	\		
94465	6	False	True
False			
100712	6	False	False
False			
153524	1	False	False
False			
85660	4	False	False
False			
100506	3	False	False
False			



	order_protocol_5	order_protocol_6	order_protocol_7
market_id_2 \			
94465	False	False	False
False			
100712	False	False	False
False			
153524	False	False	False
True			
85660	True	False	False
True			
100506	True	False	False
True			

	market_id_3	market_id_4	market_id_5	market_id_6 \
94465	False	True	False	False
100712	True	False	False	False
153524	False	False	False	False
85660	False	False	False	False
100506	False	False	False	False

	store_primary_category_13	store_primary_category_20 \
94465	False	True
100712	False	False
153524	False	False
85660	False	False
100506	False	False

	store_primary_category_24	store_primary_category_28 \
94465	False	False
100712	False	False
153524	False	False
85660	False	False
100506	False	False

	store_primary_category_38	store_primary_category_39 \
94465	False	False
100712	False	False
153524	False	False
85660	False	False
100506	False	False

	store_primary_category_46	store_primary_category_55 \
94465	False	False
100712	True	False
153524	True	False
85660	False	False
100506	False	False

store\_primary\_category\_58 store\_primary\_category\_0ther

94465	False	False
100712	False	False
153524	False	False
85660	False	False
100506	True	False

### 3. Exploratory Data Analysis on Training Data [20 marks]

1. Analyzing the correlation between variables to identify patterns and relationships
2. Identifying and addressing outliers to ensure the integrity of the analysis
3. Exploring the relationships between variables and examining the distribution of the data for better insights

#### 3.1 Feature Distributions [7 marks]

```
df.columns.to_list()

['total_items',
 'subtotal',
 'num_distinct_items',
 'min_item_price',
 'max_item_price',
 'total_onshift_dashers',
 'total_busy_dashers',
 'total_outstanding_orders',
 'distance',
 'mints_taken',
 'created_hour',
 'is_weekend',
 'day_of_week',
 'order_protocol_2',
 'order_protocol_3',
 'order_protocol_4',
 'order_protocol_5',
 'order_protocol_6',
 'order_protocol_7',
 'market_id_2',
 'market_id_3',
 'market_id_4',
 'market_id_5',
 'market_id_6',
 'store_primary_category_13',
 'store_primary_category_20',
 'store_primary_category_24',
 'store_primary_category_28',
 'store_primary_category_38',
 'store_primary_category_39',
 'store_primary_category_46',
 'store_primary_category_55',
```

```

'store_primary_category_58',
'store_primary_category_0ther']

# Define numerical and categorical columns for easy EDA and data
manipulation
numerical =
['total_items', 'subtotal', 'num_distinct_items', 'min_item_price', 'max_i
tem_price', 'total_onshift_dashers',

'total_busy_dashers', 'total_outstanding_orders', 'distance']

categorical = ['is_weekend', 'day_of_week', 'created_hour']

numerical

['total_items',
'subtotal',
'num_distinct_items',
'min_item_price',
'max_item_price',
'total_onshift_dashers',
'total_busy_dashers',
'total_outstanding_orders',
'distance']

```

### 3.1.1 [3 marks]

Plot distributions for numerical columns in the training set to understand their spread and any skewness

```

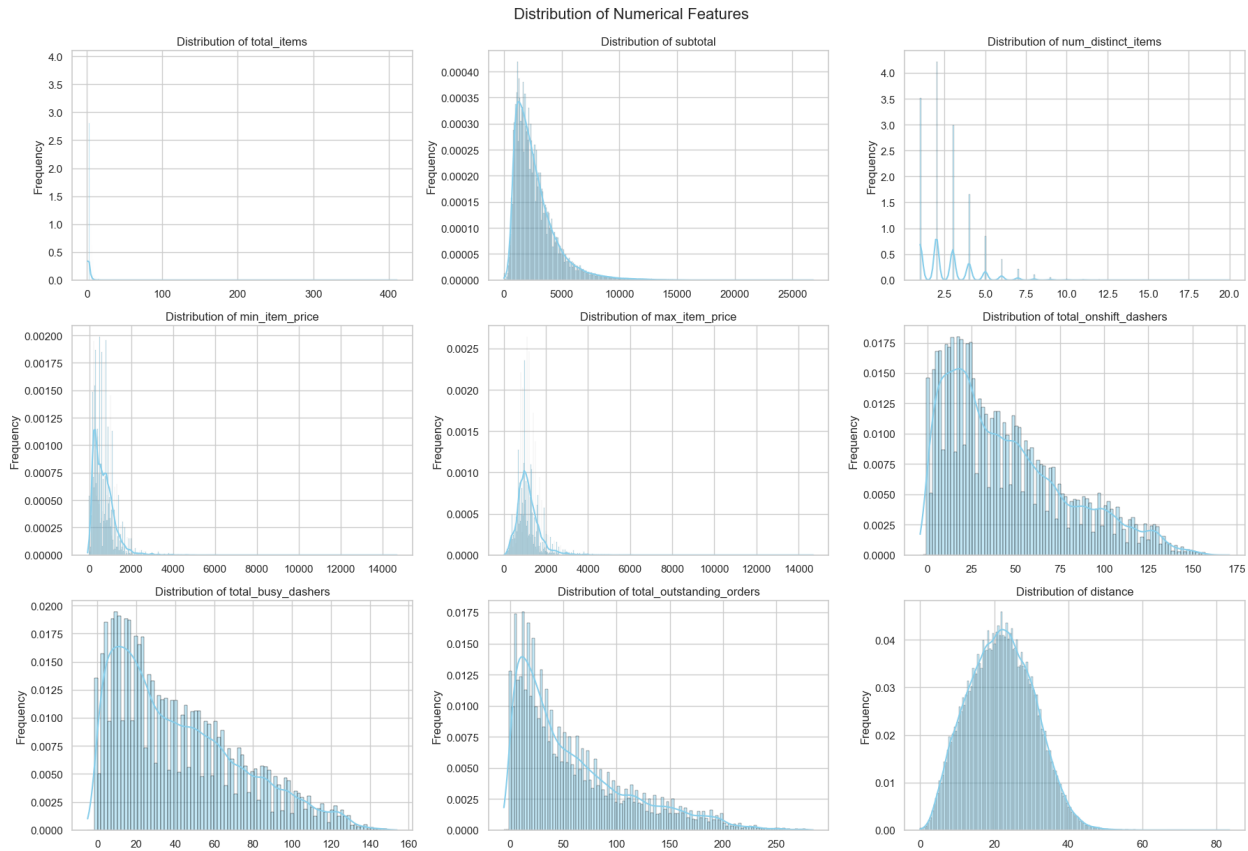
# Plot distributions for all numerical columns

# Create subplots (3 rows x 3 columns)
fig, axes = plt.subplots(nrows=3, ncols=3, figsize=(18, 12))
axes = axes.flatten() # flatten 2D axes to 1D list

# Plot each histogram
for i, col in enumerate(numerical):
    sns.histplot(df[col], kde=True, ax=axes[i], color='skyblue',
edgecolor='black', stat='density')
    axes[i].set_title(f'Distribution of {col}', fontsize=12)
    axes[i].set_xlabel('')
    axes[i].set_ylabel('Frequency')

# Adjust layout
plt.tight_layout()
plt.suptitle('Distribution of Numerical Features', fontsize=16,
y=1.02)
plt.show()

```



### 3.1.2 [2 marks]

Check the distribution of categorical features

```
df.head()
```

	total_items	subtotal	num_distinct_items	min_item_price
0	4	3441	4	557
1	1	1900	1	1400
2	4	4771	3	820
3	1	1525	1	1525
4	2	3620	2	1425

	total_onshift_dashers	total_busy_dashers	total_outstanding_orders
0	33	14	21
1	1	2	2

2	8	6	18
3	5	6	8
4	5	5	7

	distance	mints_taken	created_hour	is_weekend	day_of_week	\
0	34.44	47.0	22	0	4	
1	27.60	44.0	21	0	1	
2	11.56	55.0	0	0	0	
3	31.80	59.0	3	0	3	
4	8.20	46.0	2	0	1	

	order_protocol_2	order_protocol_3	order_protocol_4	order_protocol_5	\
0	False	False	False	False	
1	True	False	False	False	
2	False	True	False	False	
3	False	False	False	False	
4	False	False	False	False	

	order_protocol_6	order_protocol_7	market_id_2	market_id_3	market_id_4	\
0	False	False	False	False	False	
1	False	False	True	False	False	
2	False	False	True	False	False	
3	False	False	False	False	False	
4	False	False	False	False	False	

	market_id_5	market_id_6	store_primary_category_13	\
0	False	False	False	
1	False	False	False	
2	False	False	False	
3	False	False	False	
4	False	False	False	

	store_primary_category_20	store_primary_category_24	\
0	False	False	
1	False	False	

2	False	False
3	False	False
4	False	False

	store_primary_category_28	store_primary_category_38	\
0	False	False	
1	False	False	
2	False	False	
3	False	True	
4	False	True	

	store_primary_category_39	store_primary_category_46	\
0	False	False	
1	False	True	
2	False	False	
3	False	False	
4	False	False	

	store_primary_category_55	store_primary_category_58	\
0	False	False	
1	False	False	
2	False	False	
3	False	False	
4	False	False	

	store_primary_category_Other
0	False
1	False
2	True
3	False
4	False

*# Distribution of categorical columns*

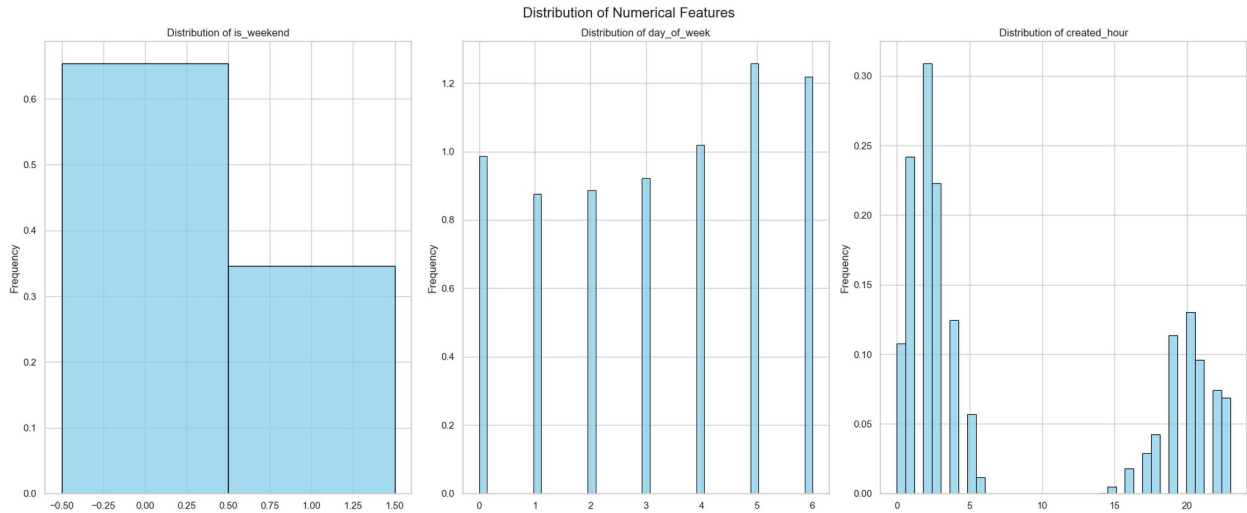
```
fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(20, 8))
axes = axes.flatten() # flatten 2D axes to 1D list
```

*# Plot each histogram*

```
for i, col in enumerate(categorical):
    sns.histplot(df[col], ax=axes[i], color='skyblue',
edgecolor='black', stat='density')
    axes[i].set_title(f'Distribution of {col}', fontsize=12)
    axes[i].set_xlabel('')
    axes[i].set_ylabel('Frequency')
```

*# Adjust layout*

```
plt.tight_layout()
plt.suptitle('Distribution of Numerical Features', fontsize=16,
y=1.02)
plt.show()
```




### 3.1.3 [2 mark]

Visualise the distribution of the target variable to understand its spread and any skewness

```
# Distribution of time_taken
```

```
plt.figure(figsize=(10, 6))
sns.histplot(y, kde=True, color='skyblue', edgecolor='black',
stat='density')
plt.show()
```



## 3.2 Relationships Between Features [3 marks]

### 3.2.1 [3 marks]

Scatter plots for important numerical and categorical features to observe how they relate to `time_taken`

```
# Scatter plot to visualise the relationship between time_taken and other features
```

```
fig, axes = plt.subplots(nrows=3, ncols=3, figsize=(15, 8))
axes = axes.flatten() # flatten 2D axes to 1D list
```

```
# Plot each histogram
```

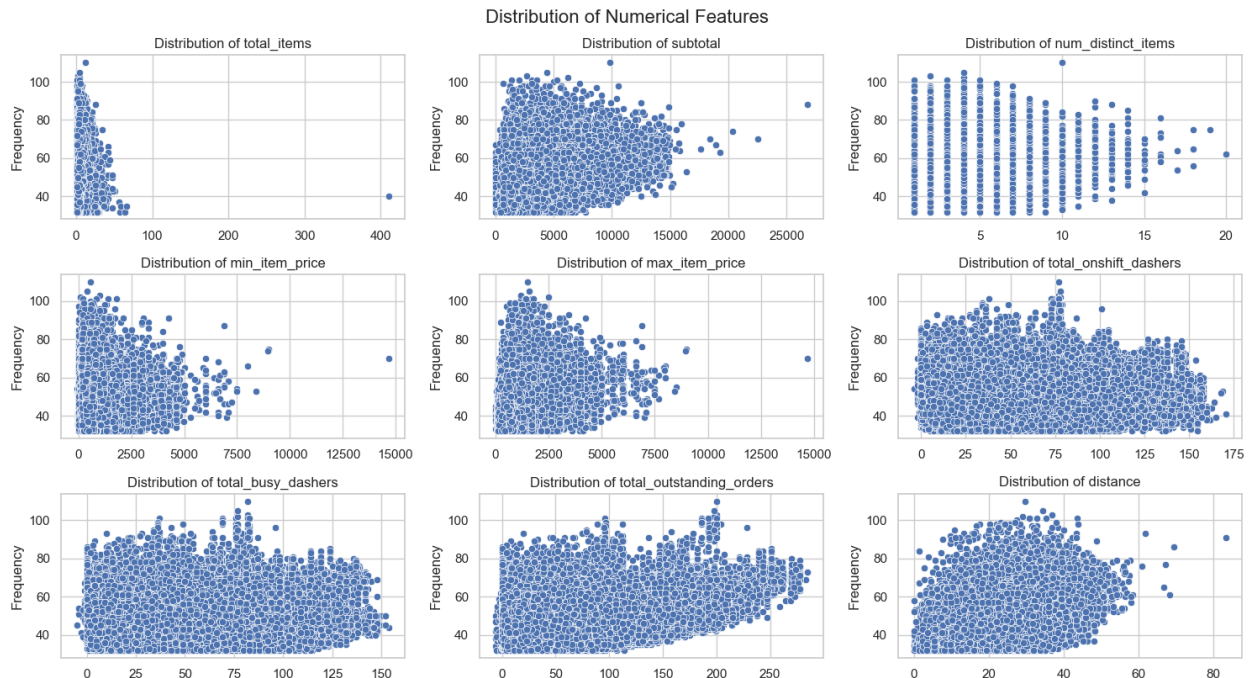
```
for i, col in enumerate(numerical):
    sns.scatterplot(x=df[col], y=y, ax=axes[i])
    axes[i].set_title(f'Distribution of {col}', fontsize=12)
    axes[i].set_xlabel('')
    axes[i].set_ylabel('Frequency')
```

```
# Adjust layout
```

```
plt.tight_layout()
```

```
plt.suptitle('Distribution of Numerical Features', fontsize=16,
y=1.02)
```

```
plt.show()
```



```
df.head()
```



	total_items	subtotal	num_distinct_items	min_item_price
max_item_price \				
0	4	3441	4	557
1239				
1	1	1900	1	1400
1400				
2	4	4771	3	820
1604				
3	1	1525	1	1525
1525				
4	2	3620	2	1425
2195				

	total_onshift_dashers	total_busy_dashers	total_outstanding_orders
\			
0	33	14	21
1	1	2	2
2	8	6	18
3	5	6	8
4	5	5	7

	distance	mints_taken	created_hour	is_weekend	day_of_week	\
0	34.44	47.0	22	0	4	
1	27.60	44.0	21	0	1	
2	11.56	55.0	0	0	0	
3	31.80	59.0	3	0	3	
4	8.20	46.0	2	0	1	

	order_protocol_2	order_protocol_3	order_protocol_4
order_protocol_5 \			
0	False	False	False
False			
1	True	False	False
False			
2	False	True	False
False			
3	False	False	False
False			
4	False	False	False
False			

	order_protocol_6	order_protocol_7	market_id_2	market_id_3
market_id_4 \				
0	False	False	False	False
False				
1	False	False	True	False

False				
2	False	False	True	False
False				
3	False	False	False	False
False				
4	False	False	False	False
False				

	market_id_5	market_id_6	store_primary_category_13	\
0	False	False	False	
1	False	False	False	
2	False	False	False	
3	False	False	False	
4	False	False	False	

	store_primary_category_20	store_primary_category_24	\
0	False	False	
1	False	False	
2	False	False	
3	False	False	
4	False	False	

	store_primary_category_28	store_primary_category_38	\
0	False	False	
1	False	False	
2	False	False	
3	False	True	
4	False	True	

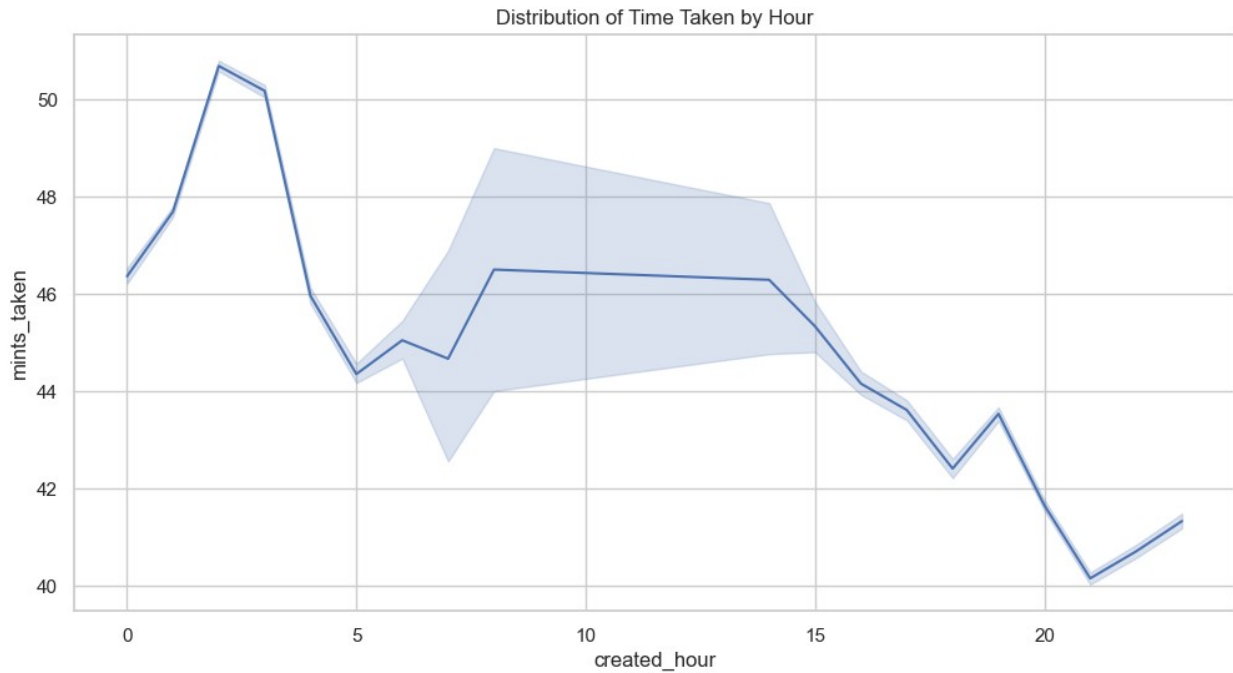
	store_primary_category_39	store_primary_category_46	\
0	False	False	
1	False	True	
2	False	False	
3	False	False	
4	False	False	

	store_primary_category_55	store_primary_category_58	\
0	False	False	
1	False	False	
2	False	False	
3	False	False	
4	False	False	


	store_primary_category_Other
0	False
1	False
2	True
3	False
4	False

```
# Show the distribution of time_taken for different hours
```

```
plt.figure(figsize=(12, 6))  
sns.lineplot(data=df, x='created_hour', y=y )  
plt.title('Distribution of Time Taken by Hour')  
plt.show()
```



```
plt.figure(figsize=(12, 6))  
sns.boxplot(data=df, x='created_hour', y=y, color='skyblue',  
fliersize=0, linewidth=1.5, width=0.5)  
plt.show()
```



### 3.3 Correlation Analysis [5 marks]

Check correlations between numerical features to identify which variables are strongly related to `time_taken`

#### 3.3.1 [3 marks]


Plot a heatmap to display correlations

```
# Plot the heatmap of the correlation matrix

corr_matrix = df.corr(numeric_only=True)
target_col = 'mints_taken'

target_corr = corr_matrix[target_col].abs().sort_values()

plt.figure(figsize=(20,20))
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap='coolwarm',
            square=True, cbar_kws={"shrink": .8})
plt.title('Correlation Matrix', fontsize=16)
plt.show()
```



target\_corr[target\_corr < 0.1]

```
order_protocol_6      0.001738
order_protocol_7      0.004744
market_id_6           0.005360
store_primary_category_24 0.009810
order_protocol_2      0.019410
min_item_price        0.022753
store_primary_category_20 0.026218
order_protocol_3      0.028757
market_id_4           0.036712
store_primary_category_Other 0.036989
store_primary_category_55 0.041729
store_primary_category_13 0.042540
market_id_5           0.042567
day_of_week           0.045878
```

```

store_primary_category_46      0.055209
store_primary_category_38      0.062337
store_primary_category_39      0.069962
store_primary_category_28      0.074015
market_id_3                    0.082070
store_primary_category_58      0.082774
order_protocol_5               0.084833
market_id_2                    0.094109
Name: mints_taken, dtype: float64

```

In correlation usually values less than 0.4 are considered weak, but that removes too much columns, we are going to remove columns which have correlated values less than 0.1.

### 3.3.2 [2 marks]

Drop the columns with weak correlations with the target variable

```

# Drop 3-5 weakly correlated columns from training dataset

weak_columns = target_corr[target_corr < 0.1].index.tolist()
print(f"Columns with weak correlation to target: {weak_columns}")
X_train = X_train.drop(columns=weak_columns)

Columns with weak correlation to target: ['order_protocol_6',
'order_protocol_7', 'market_id_6', 'store_primary_category_24',
'order_protocol_2', 'min_item_price', 'store_primary_category_20',
'order_protocol_3', 'market_id_4', 'store_primary_category_0ther',
'store_primary_category_55', 'store_primary_category_13',
'market_id_5', 'day_of_week', 'store_primary_category_46',
'store_primary_category_38', 'store_primary_category_39',
'store_primary_category_28', 'market_id_3',
'store_primary_category_58', 'order_protocol_5', 'market_id_2']

X_train.columns.tolist()

['total_items',
'subtotal',
'num_distinct_items',
'max_item_price',
'total_onshift_dashers',
'total_busy_dashers',
'total_outstanding_orders',
'distance',
'created_hour',
'is_weekend',
'order_protocol_4']

```

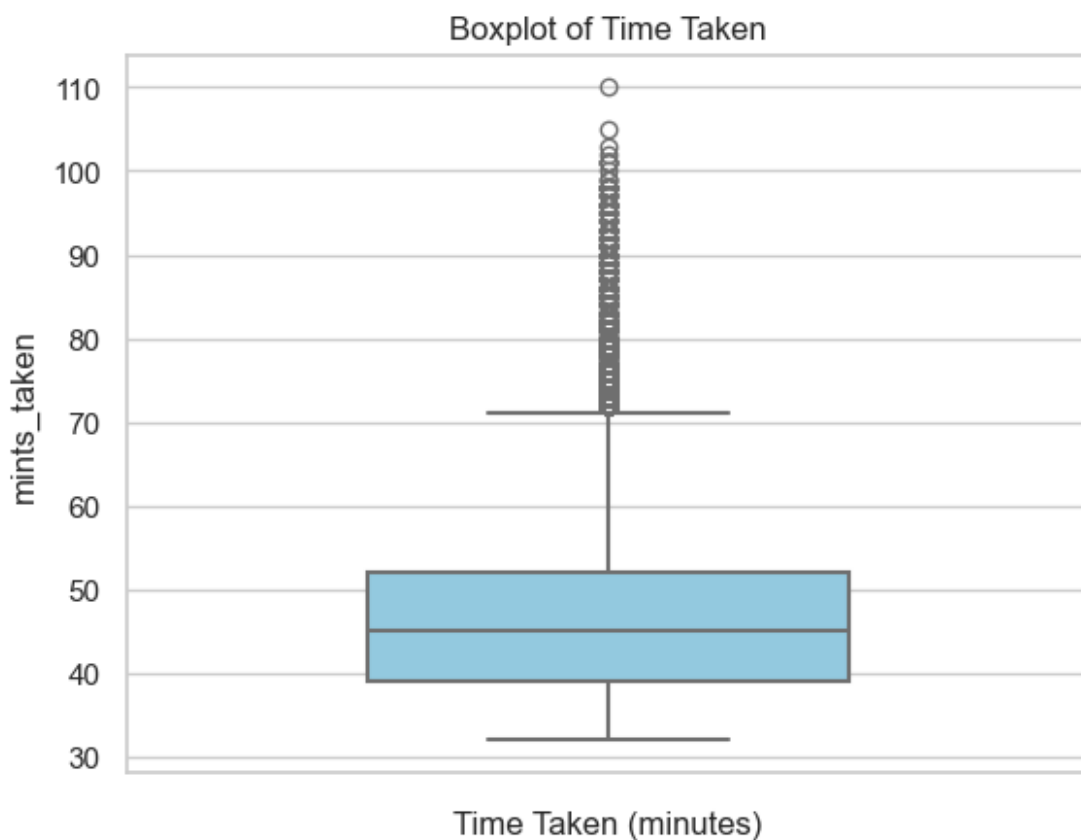
### 3.4 Handling the Outliers [5 marks]

#### 3.4.1 [2 marks]

Visualise potential outliers for the target variable and other numerical features using boxplots

```
# Boxplot for time_taken
```

```
sns.boxplot(y=df['mints_taken'], color='skyblue', linewidth=1.5, width=0.5)
plt.title('Boxplot of Time Taken')
plt.xlabel('Time Taken (minutes)')
plt.show()
```



#### 3.4.2 [3 marks]

Handle outliers present in all columns

```
# Handle outliers
```

```
numeric_cols = df.select_dtypes(include='number').columns
```

```
# Set number of plots
```

```
n_cols = 3
```

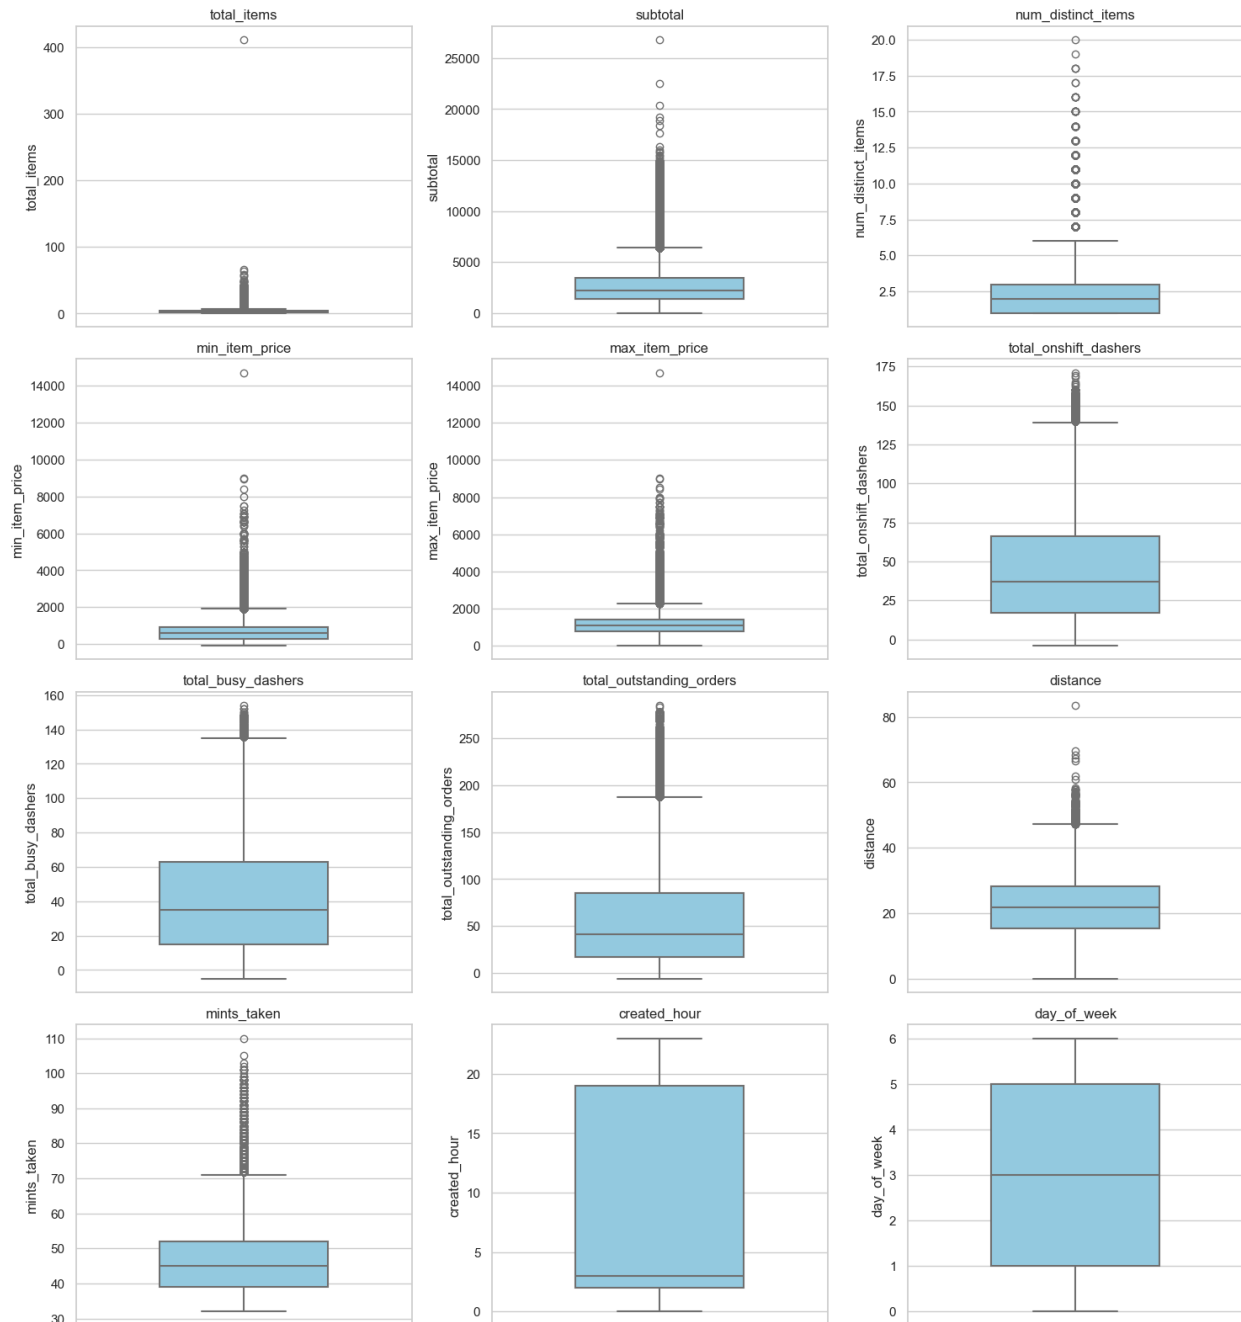
```
n_rows = -(-len(numeric_cols) // n_cols) # Ceiling division

# Set plot size
plt.figure(figsize=(5 * n_cols, 4 * n_rows))

for idx, col in enumerate(numeric_cols, 1):
    plt.subplot(n_rows, n_cols, idx)
    sns.boxplot(y=df[col], color='skyblue', linewidth=1.5, width=0.5)
    plt.title(col)

plt.tight_layout()
plt.show()
```





```
def remove_outliers(df, columns, factor=3):
    df_clean = df.copy()
    for col in columns:
        Q1 = df_clean[col].quantile(0.25)
        Q3 = df_clean[col].quantile(0.75)
        IQR = Q3 - Q1
        lower = Q1 - factor * IQR
        upper = Q3 + factor * IQR
        df_clean = df_clean[(df_clean[col] >= lower) & (df_clean[col]
<= upper)]
```

```
    return df_clean

numeric_cols = df.select_dtypes(include='number').columns.tolist()
df_cleaned = remove_outliers(df, numeric_cols, factor=3)
df_cleaned.shape

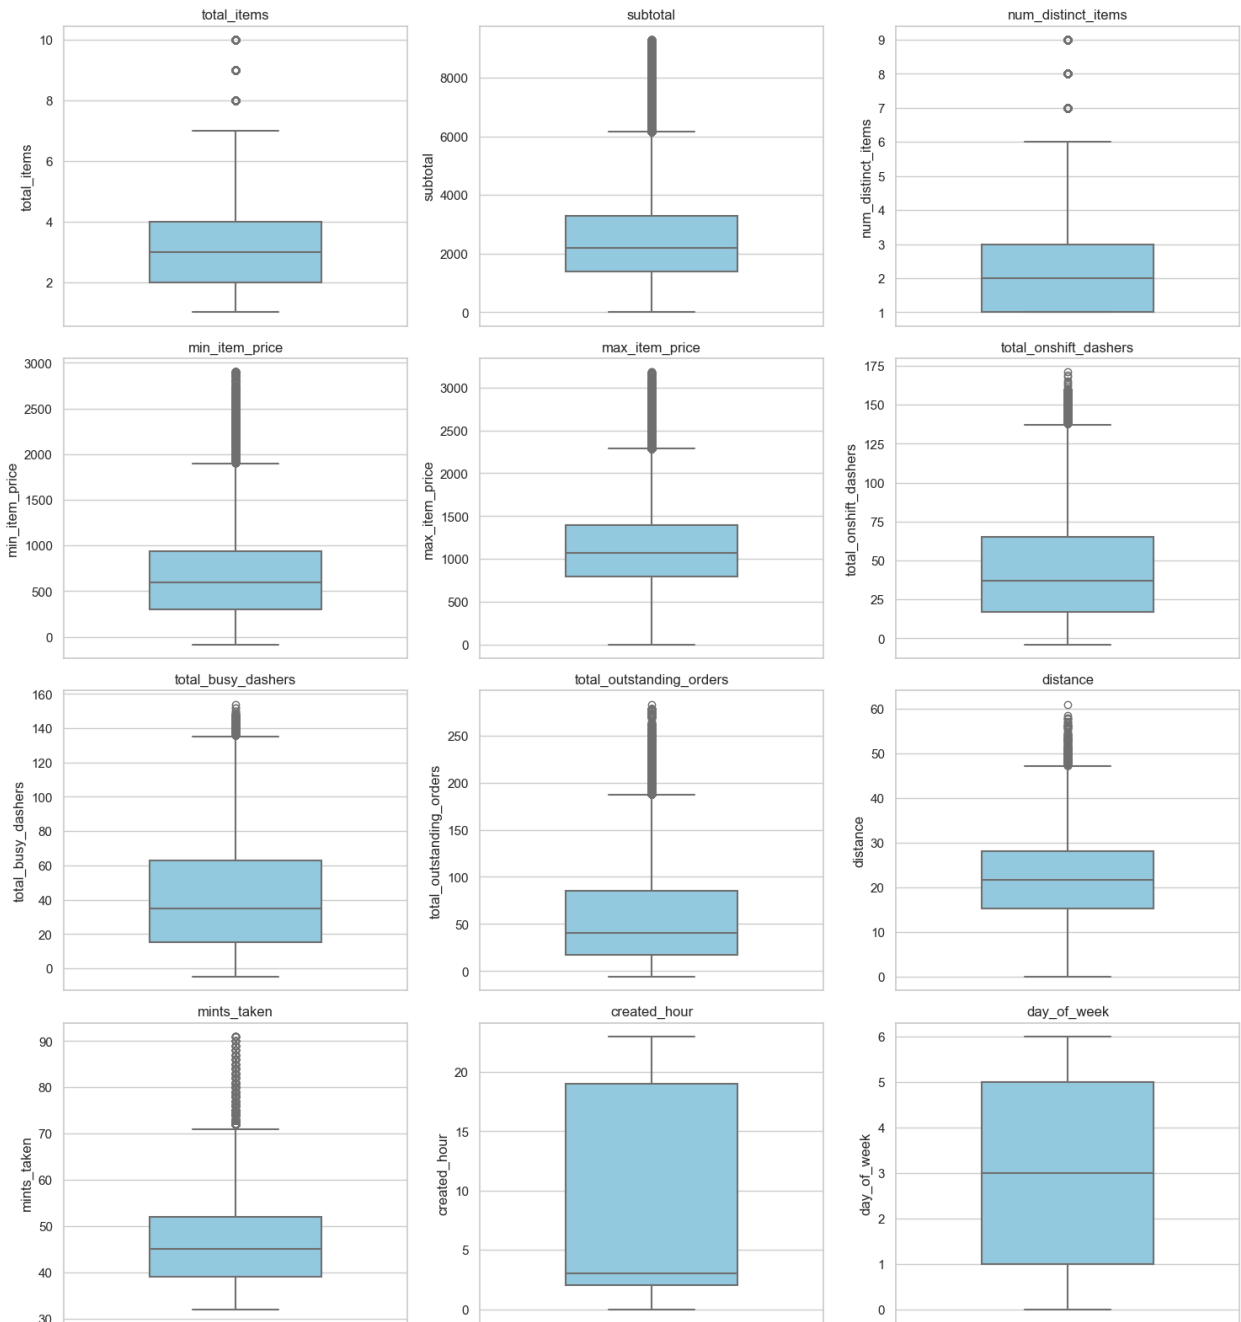
(169730, 34)

# Set number of plots
n_cols = 3
n_rows = -(-len(numeric_cols) // n_cols) # Ceiling division

# Set plot size
plt.figure(figsize=(5 * n_cols, 4 * n_rows))

for idx, col in enumerate(numeric_cols, 1):
    plt.subplot(n_rows, n_cols, idx)
    sns.boxplot(y=df_cleaned[col], color='skyblue', linewidth=1.5,
width=0.5)
    plt.title(col)

plt.tight_layout()
plt.show()
```



## 4. Exploratory Data Analysis on Validation Data [optional]

Optionally, perform EDA on test data to see if the distribution match with the training data

```
# Define numerical and categorical columns for easy EDA and data manipulation
```

## 4.1 Feature Distributions

### 4.1.1

Plot distributions for numerical columns in the validation set to understand their spread and any skewness

```
# Plot distributions for all numerical columns
```

### 4.1.2

Check the distribution of categorical features

```
# Distribution of categorical columns
```

### 4.1.3

Visualise the distribution of the target variable to understand its spread and any skewness

```
# Distribution of time_taken
```

## 4.2 Relationships Between Features

Scatter plots for numerical features to observe how they relate to each other, especially to `time_taken`

```
# Scatter plot to visualise the relationship between time_taken and  
other features
```

## 4.3 Drop the columns with weak correlations with the target variable

```
# Drop the weakly correlated columns from training dataset
```

# 5. Model Building [15 marks]

## Import Necessary Libraries

```
# Import libraries  
import scipy.stats as stats  
import statsmodels.api as sm  
from sklearn.feature_selection import RFE  
from sklearn.linear_model import LinearRegression as LR  
from sklearn.metrics import mean_squared_error, r2_score # Import  
evaluation metrics  
from sklearn.preprocessing import StandardScaler
```

## 5.1 Feature Scaling [3 marks]

```
df_cleaned[numeric_cols].head()
```

	total_items	subtotal	num_distinct_items	min_item_price
0	4	3441	4	557
1	1	1900	1	1400
2	4	4771	3	820
3	1	1525	1	1525
4	2	3620	2	1425

	total_onshift_dashers	total_busy_dashers	total_outstanding_orders
0	33	14	21
1	1	2	2
2	8	6	18
3	5	6	8
4	5	5	7

	distance	mints_taken	created_hour	day_of_week
0	34.44	47.0	22	4
1	27.60	44.0	21	1
2	11.56	55.0	0	0
3	31.80	59.0	3	3
4	8.20	46.0	2	1

```
numeric_cols
```

```
['total_items',  
 'subtotal',  
 'num_distinct_items',  
 'min_item_price',  
 'max_item_price',  
 'total_onshift_dashers',  
 'total_busy_dashers',  
 'total_outstanding_orders',  
 'distance',  
 'mints_taken',  
 'created_hour',  
 'day_of_week']
```

```

print(X_train.columns.to_list())
print(X_test.columns.to_list())

['total_items', 'subtotal', 'num_distinct_items', 'max_item_price',
'total_onshift_dashers', 'total_busy_dashers',
'total_outstanding_orders', 'distance', 'created_hour', 'is_weekend',
'order_protocol_4']
['total_items', 'subtotal', 'num_distinct_items', 'min_item_price',
'max_item_price', 'total_onshift_dashers', 'total_busy_dashers',
'total_outstanding_orders', 'distance', 'created_hour', 'is_weekend',
'day_of_week', 'order_protocol_2', 'order_protocol_3',
'order_protocol_4', 'order_protocol_5', 'order_protocol_6',
'order_protocol_7', 'market_id_2', 'market_id_3', 'market_id_4',
'market_id_5', 'market_id_6', 'store_primary_category_13',
'store_primary_category_20', 'store_primary_category_24',
'store_primary_category_28', 'store_primary_category_38',
'store_primary_category_39', 'store_primary_category_46',
'store_primary_category_55', 'store_primary_category_58',
'store_primary_category_other']

X_train['is_weekend'] = X_train['is_weekend'].astype(int)
X_train['order_protocol_4'] = X_train['order_protocol_4'].astype(int)

# Apply scaling to the numerical columns
X_test = X_test[X_train.columns]
num_cols = X_train.select_dtypes(include='number').columns.to_list()
scaler = StandardScaler()

X_train_unscaled = X_train.copy()
X_test_unscaled = X_test.copy()

X_train[num_cols] = scaler.fit_transform(X_train[num_cols])
X_test[num_cols] = scaler.transform(X_test[num_cols])

# X_train = pd.DataFrame(scaler.fit_transform(X_train),
# columns=num_cols, index=X_train.index)
# X_test = pd.DataFrame(scaler.transform(X_test), columns=num_cols,
# index=X_test.index)

num_cols

['total_items',
'subtotal',
'num_distinct_items',
'max_item_price',
'total_onshift_dashers',
'total_busy_dashers',
'total_outstanding_orders',
'distance',
'created_hour',

```

```
'is_weekend',  
'order_protocol_4']
```

```
X_train.head()
```

	total_items	subtotal	num_distinct_items	max_item_price	\
94465	-0.075210	0.920885	0.198861	0.420814	
100712	0.287056	0.040812	0.198861	0.338664	
153524	1.373857	-0.564273	1.427533	-1.036469	
85660	-0.075210	-0.665757	0.198861	-1.036469	
100506	-0.799744	-0.852902	-1.029812	-0.775729	

	total_onshift_dashers	total_busy_dashers	
total_outstanding_orders	\		
94465	1.130856	0.656539	
0.810319			
100712	-0.663302	-0.804025	-
0.857394			
153524	-1.242063	-1.270162	-
1.103761			
85660	-0.113480	0.003947	-
0.326758			
100506	0.436343	0.594387	
0.393391			

	distance	created_hour	is_weekend	order_protocol_4
94465	0.558203	-0.747442	1.376019	-0.330386
100712	0.118915	-0.632199	1.376019	-0.330386
153524	-0.919819	-0.401714	-0.726734	-0.330386
85660	-1.496384	-0.632199	-0.726734	-0.330386
100506	-1.079976	1.326924	-0.726734	-0.330386

```
# Create/Initialise the model
```

```
lr = LR()  
lr.fit(X_train, y_train)
```

```
LinearRegression()
```

```
X_train.dtypes
```

total_items	float64
subtotal	float64
num_distinct_items	float64
max_item_price	float64
total_onshift_dashers	float64
total_busy_dashers	float64
total_outstanding_orders	float64
distance	float64
created_hour	float64
is_weekend	float64

```
order_protocol_4          float64
dtype: object
```

```
X_train.head()
```

	total_items	subtotal	num_distinct_items	max_item_price	\
94465	-0.075210	0.920885	0.198861	0.420814	
100712	0.287056	0.040812	0.198861	0.338664	
153524	1.373857	-0.564273	1.427533	-1.036469	
85660	-0.075210	-0.665757	0.198861	-1.036469	
100506	-0.799744	-0.852902	-1.029812	-0.775729	

	total_onshift_dashers	total_busy_dashers	
total_outstanding_orders			\
94465	1.130856	0.656539	
0.810319			
100712	-0.663302	-0.804025	-
0.857394			
153524	-1.242063	-1.270162	-
1.103761			
85660	-0.113480	0.003947	-
0.326758			
100506	0.436343	0.594387	
0.393391			

	distance	created_hour	is_weekend	order_protocol_4
94465	0.558203	-0.747442	1.376019	-0.330386
100712	0.118915	-0.632199	1.376019	-0.330386
153524	-0.919819	-0.401714	-0.726734	-0.330386
85660	-1.496384	-0.632199	-0.726734	-0.330386
100506	-1.079976	1.326924	-0.726734	-0.330386

```
X_train['is_weekend'] = X_train['is_weekend'].astype(int)
X_train['order_protocol_4'] = X_train['order_protocol_4'].astype(int)
```

```
# Train the model using the training data
```

```
X_train = sm.add_constant(X_train)
model = sm.OLS(y_train, X_train).fit()
model.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

#### OLS Regression Results

```
=====
=====
Dep. Variable:          mints_taken    R-squared:
0.862
Model:                  OLS    Adj. R-squared:
0.862
Method:                 Least Squares    F-statistic:
```



Date: Sat, 28 Jun 2025 Prob (F-statistic): 0.00

3.2718e+05

6.544e+05

6.545e+05

Covariance Type: nonrobust

	coef	std err	t	P> t
--	------	---------	---	------

\_\_\_\_\_

const		45.8191	0.015	5020.129	0.000
45.794	45.844				

total_items	011112	01019	71999	01000
-0.144	-0.085			

2.270	2.342
-------	-------

0.850	0.919
-------	-------

0.413	0.466
-------	-------

-12.795	-12.662
---------	---------

-4.704	-4.575				
total	total	total	total	total	total
10.3501	0.001	501.064	0.000		

18.297	18.419				
distance		4.1576	0.010	431.050	0.000

4.138	4.177				
created	hour	-2.3426	0.011	-207.082	0.000

-2.204	-2.221				
is_weekend		1.2660	0.021	60.402	0.000

order protocol 4	-0.1989	0.012	-17.250	0.000
------------------	---------	-------	---------	-------

=====

Omnibus: 36731.307 Durbin-Watson:

Prob(Omnibus): 0.000 Jarque-Bera (JB):

144255.998

Skew:	1.448	Prob(JB):
0.00		
Kurtosis:	7.444	Cond. No.
7.68		

=====

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

*# Make predictions*

X\_test\_const = sm.add\_constant(X\_test, has\_constant='add')

y\_test\_pred = model.predict(X\_test\_const)

*# y\_test\_pred*

*# Find results for evaluation metrics*

mse = mean\_squared\_error(y\_test, y\_test\_pred)

rmse = mse \*\* 0.5

r2 = r2\_score(y\_test, y\_test\_pred)

print(f'The MSE is {mse}, \n RMSE is {rmse},\n R2\_score is {r2}')

The MSE is 12.614605658256282,

RMSE is 3.5517046130353074,

R2\_score is 0.8555764888025942

print(y\_test.dtypes)

print(y\_test\_pred.dtypes)

float64

float64

y\_test\_pred = pd.Series(y\_test\_pred).astype('float64')

sns.regplot(x=y\_test, y=y\_test\_pred, scatter\_kws={"alpha":0.5,"color":  
"red"})

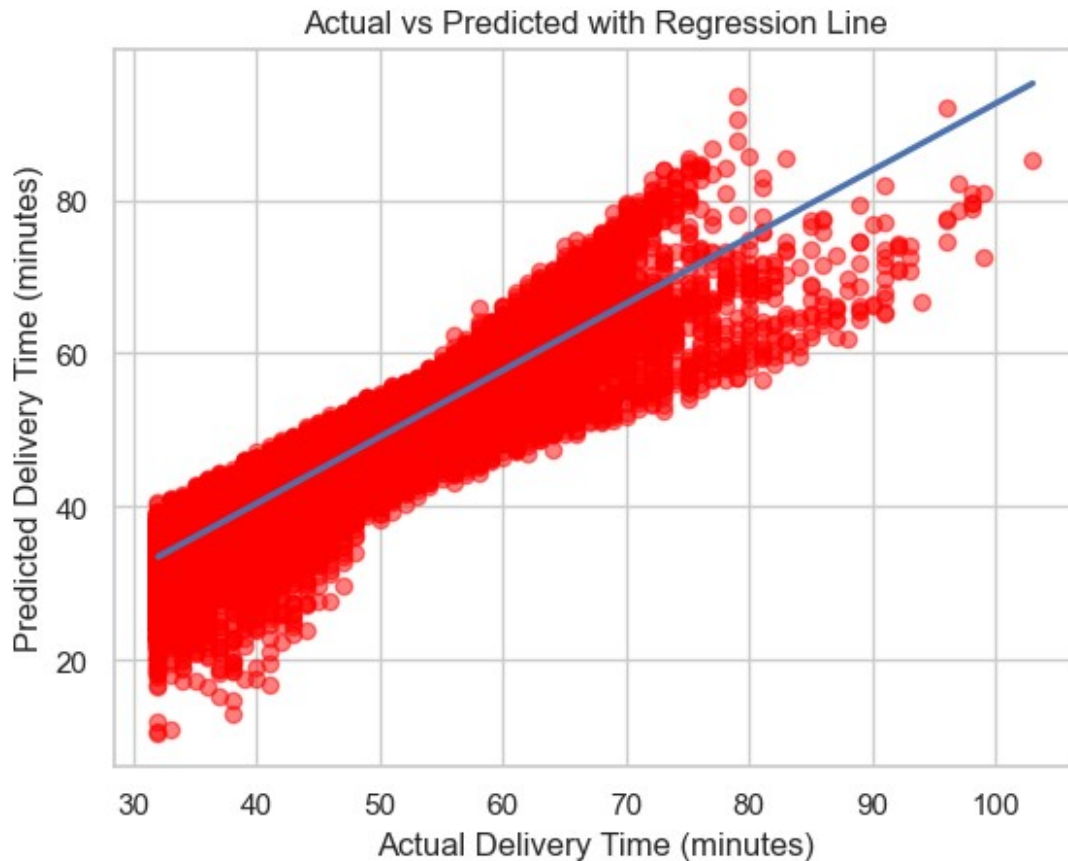
plt.xlabel("Actual Delivery Time (minutes)")

plt.ylabel("Predicted Delivery Time (minutes)")

plt.title("Actual vs Predicted with Regression Line")

plt.grid(True)

plt.show()



Note that we have 12 (depending on how you select features) training features. However, not all of them would be useful. Let's say we want to take the most relevant 8 features.

We will use Recursive Feature Elimination (RFE) here.

For this, you can look at the coefficients / p-values of features from the model summary and perform feature elimination, or you can use the RFE module provided with *scikit-learn*.

### 5.3 Build the model and fit RFE to select the most important features [7 marks]

For RFE, we will start with all features and use the RFE method to recursively reduce the number of features one-by-one.

After analysing the results of these iterations, we select the one that has a good balance between performance and number of features.

```
# Loop through the number of features and test the model
res = []
for i in range(1, X_test.shape[1]+1):

    rfe = RFE(lr, n_features_to_select=i)
    rfe = rfe.fit(X_train,y_train)

    features = X_train.columns[rfe.support_]
```

```

lr.fit(X_train[features], y_train)

y_pred = lr.predict(X_train[features])

mse = mean_squared_error(y_pred=y_pred, y_true=y_train)
rmse = mse ** 0.5
r2 = r2_score(y_pred=y_pred, y_true=y_train)

res.append([i,mse, rmse, r2])

res_df = pd.DataFrame(res, columns=['features', 'MSE', 'RMSE', 'R2'])
print(res_df)

```

	features	MSE	RMSE	R2
0	1	73.950613	8.599454	0.148552
1	2	48.815698	6.986823	0.437949
2	3	46.852122	6.844861	0.460557
3	4	28.479874	5.336654	0.672090
4	5	16.880073	4.108537	0.805647
5	6	12.637922	3.554986	0.854490
6	7	12.263652	3.501950	0.858799
7	8	12.116064	3.480814	0.860499
8	9	11.982722	3.461607	0.862034
9	10	11.949522	3.456808	0.862416
10	11	11.944010	3.456011	0.862480

```

rfe = RFE(lr, n_features_to_select=10)
rfe = rfe.fit(X_train,y_train)
list(zip(X_train.columns, rfe.support_, rfe.ranking_))

```

```

[('const', False, 3),
 ('total_items', False, 2),
 ('subtotal', True, 1),
 ('num_distinct_items', True, 1),
 ('max_item_price', True, 1),
 ('total_onshift_dashers', True, 1),
 ('total_busy_dashers', True, 1),
 ('total_outstanding_orders', True, 1),
 ('distance', True, 1),
 ('created_hour', True, 1),
 ('is_weekend', True, 1),
 ('order_protocol_4', True, 1)]

```

```

X_train.columns.to_list()

```

```

['const',
 'total_items',
 'subtotal',
 'num_distinct_items',
 'max_item_price',

```

```

'total_onshift_dashers',
'total_busy_dashers',
'total_outstanding_orders',
'distance',
'created_hour',
'is_weekend',
'order_protocol_4']

# Build the final model with selected number of features

X_train_new = X_train[['subtotal', 'num_distinct_items',
'max_item_price', 'total_onshift_dashers', 'total_busy_dashers',
'total_outstanding_orders', 'distance', 'created_hour', 'is_weekend',
'order_protocol_4']]
X_train_new = sm.add_constant(X_train_new)
model = sm.OLS(y_train,X_train_new).fit()
model.summary()

<class 'statsmodels.iolib.summary.Summary'>
"""
                                OLS Regression Results

=====
=====
Dep. Variable:                  mints_taken    R-squared:
0.862
Model:                            OLS        Adj. R-squared:
0.862
Method:                  Least Squares    F-statistic:
7.712e+04
Date:                  Sat, 28 Jun 2025    Prob (F-statistic):
0.00
Time:                  06:55:12    Log-Likelihood:    -
3.2721e+05
No. Observations:                  123043    AIC:
6.544e+05
Df Residuals:                  123032    BIC:
6.545e+05
Df Model:                            10

Covariance Type:                  nonrobust

=====
=====
                                coef    std err          t      P>|t|
[0.025    0.975]
-----
const                45.8227      0.013   3622.129      0.000
45.798    45.847

```

subtotal	2.2748	0.018	127.186	0.000
2.240      2.310				
num_distinct_items	0.8207	0.015	53.329	0.000
0.791      0.851				
max_item_price	0.4621	0.013	35.061	0.000
0.436      0.488				
total_onshift_dashers	-12.7269	0.034	-373.612	0.000
-12.794      -12.660				
total_busy_dashers	-4.6410	0.033	-140.944	0.000
-4.705      -4.576				
total_outstanding_orders	18.3578	0.031	591.721	0.000
18.297      18.419				
distance	4.1585	0.010	421.079	0.000
4.139      4.178				
created_hour	-2.2432	0.011	-207.098	0.000
-2.264      -2.222				
is_weekend	1.2661	0.021	60.393	0.000
1.225      1.307				
order_protocol_4	-0.2111	0.011	-18.489	0.000
-0.233      -0.189				

```
=====
=====
Omnibus:                 36731.693   Durbin-Watson:
2.001
Prob(Omnibus):           0.000   Jarque-Bera (JB):
144302.228
Skew:                    1.448   Prob(JB):
0.00
Kurtosis:                7.445   Cond. No.
7.64
=====
=====
```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
"""
```

```
X_test_new = X_test[['subtotal', 'num_distinct_items',
' max_item_price', 'total_onshift_dashers', 'total_busy_dashers',
' total_outstanding_orders', 'distance', 'created_hour', 'is_weekend',
' order_protocol_4']]
```

```
X_test_const = sm.add_constant(X_test_new, has_constant='add')
```

```
y_test_pred = model.predict(X_test_const)
```

```
residual = y_test - y_test_pred
```

## 6. Results and Inference [5 marks]

### 6.1 Perform Residual Analysis [3 marks]


```
print(y_test.dtypes)
print(y_test_pred.dtypes)
print(residual.dtypes)

float64
float64
float64

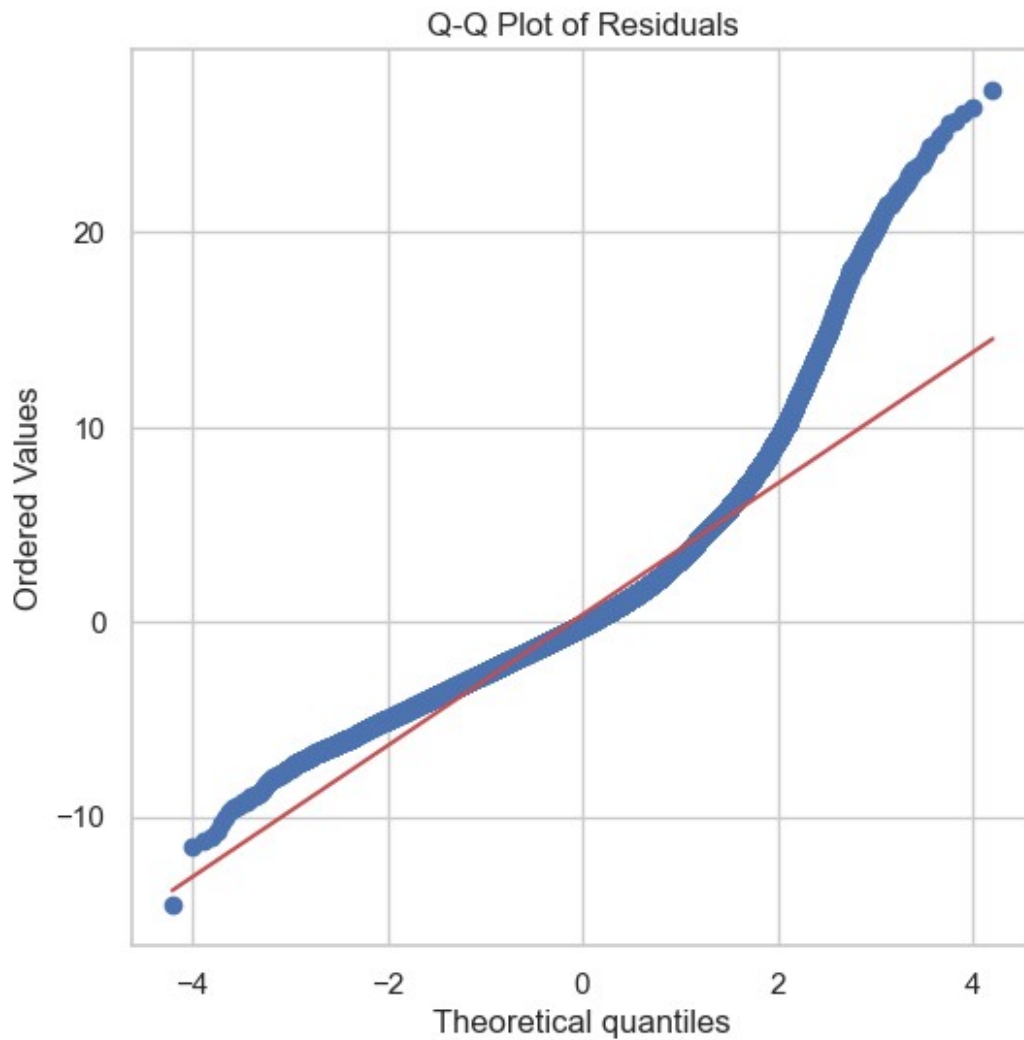
y_test_pred = pd.Series(y_test_pred).astype('float64')
residual = pd.Series(residual).astype('float64')

# Perform residual analysis using plots like residuals vs predicted
values, Q-Q plot and residual histogram

sns.regplot(x=residual, y=y_test_pred, line_kws={"alpha":0.5,"color":
"red"})
plt.xlabel("Predicted values")
plt.ylabel("Residuals")
plt.title("Residuals vs Predicted")
plt.show()
```

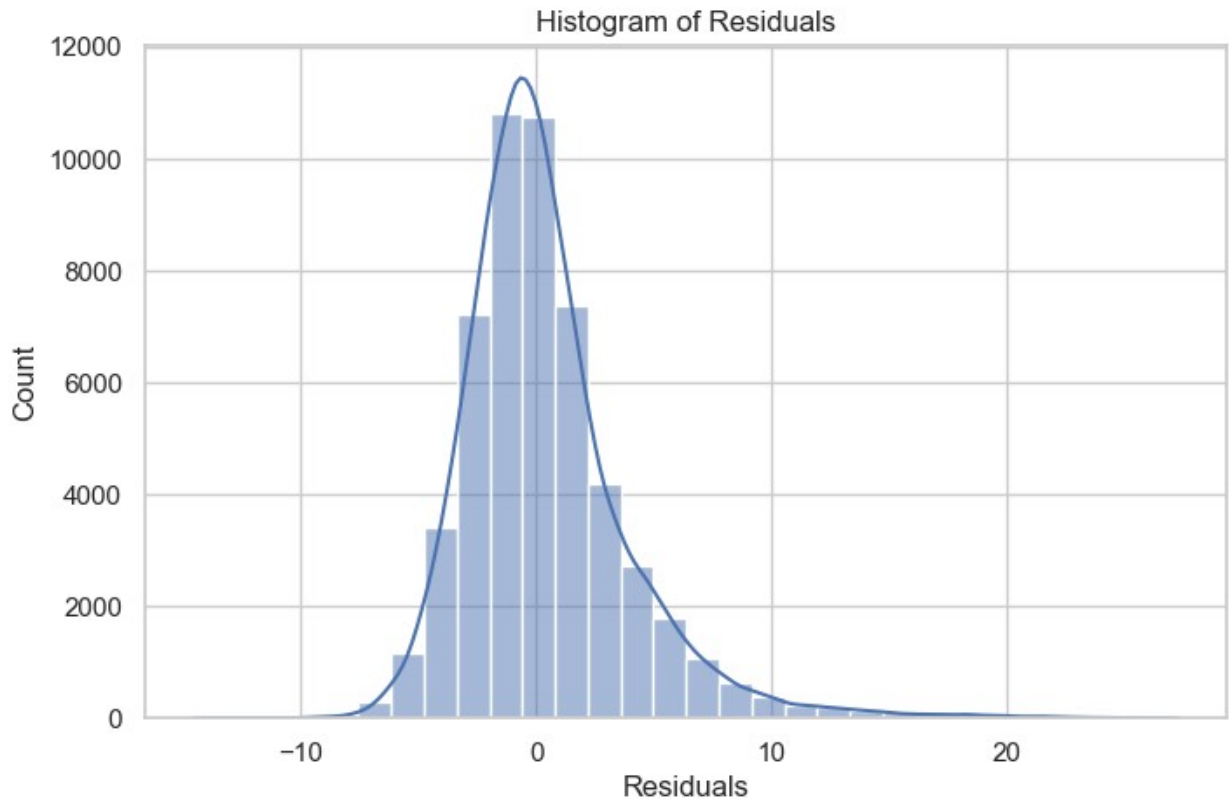


```
plt.figure(figsize=(6, 6))
stats.probplot(residual, dist="norm", plot=plt)
plt.title("Q-Q Plot of Residuals")
plt.show()
```



```
plt.figure(figsize=(8, 5))
sns.histplot(residual, kde=True, bins=30)
plt.xlabel("Residuals")
plt.title("Histogram of Residuals")
plt.show()
```





[Your inferences here:]

## 6.2 Perform Coefficient Analysis [2 marks]

Perform coefficient analysis to find how changes in features affect the target. Also, the features were scaled, so interpret the scaled and unscaled coefficients to understand the impact of feature changes on delivery time.

```
num_cols =
X_train_new.select_dtypes(include='number').columns.to_list()

# Compare the scaled vs unscaled features used in the final model

coef_scaled = model.params[num_cols].values
intercept_scaled = model.params[0]

# print(model.params[num_cols])
# print(list(zip(scaler.get_feature_names_out(), scaler.scale_)))

# To unscale: divide by standard scaler's scale
coef_unscaled = (coef_scaled[1:]) / (scaler.scale_[1:])
intercept_unscaled = intercept_scaled - np.sum(coef_scaled[1:] *
scaler.mean_[1:] / scaler.scale_[1:])
```

C:\Users\user\AppData\Local\Temp\ipykernel\_17484\2993582975.py:4:  
FutureWarning: Series.\_\_getitem\_\_ treating keys as positions is

deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use ``ser.iloc[pos]``

```
intercept_scaled = model.params[0]

coef_df = pd.DataFrame({
    "Feature": X_train_new.columns[1:],
    "Scaled Coef": coef_scaled[1:],
    "Unscaled Coef": coef_unscaled,
    "Mean": scaler.mean_[1:],
    "Std Dev": scaler.scale_[1:]
})

print(coef_df.round(4))
```

	Feature	Scaled Coef	Unscaled Coef	Mean
Std Dev				
0	subtotal	2.2748	0.0012	2697.2000
1832.8014				
1	num_distinct_items	0.8207	0.5042	2.6763
1.6278				
2	max_item_price	0.4621	0.0008	1159.3666
559.9461				
3	total_onshift_dashers	-12.7269	-0.3683	44.9215
34.5566				
4	total_busy_dashers	-4.6410	-0.1442	41.8730
32.1794				
5	total_outstanding_orders	18.3578	0.3479	58.2420
52.7669				
6	distance	4.1585	0.4757	21.8405
8.7414				
7	created_hour	-2.2432	-0.2585	8.4858
8.6774				
8	is_weekend	1.2661	2.6624	0.3456
0.4756				
9	order_protocol_4	-0.2111	-0.7086	0.0984
0.2979				

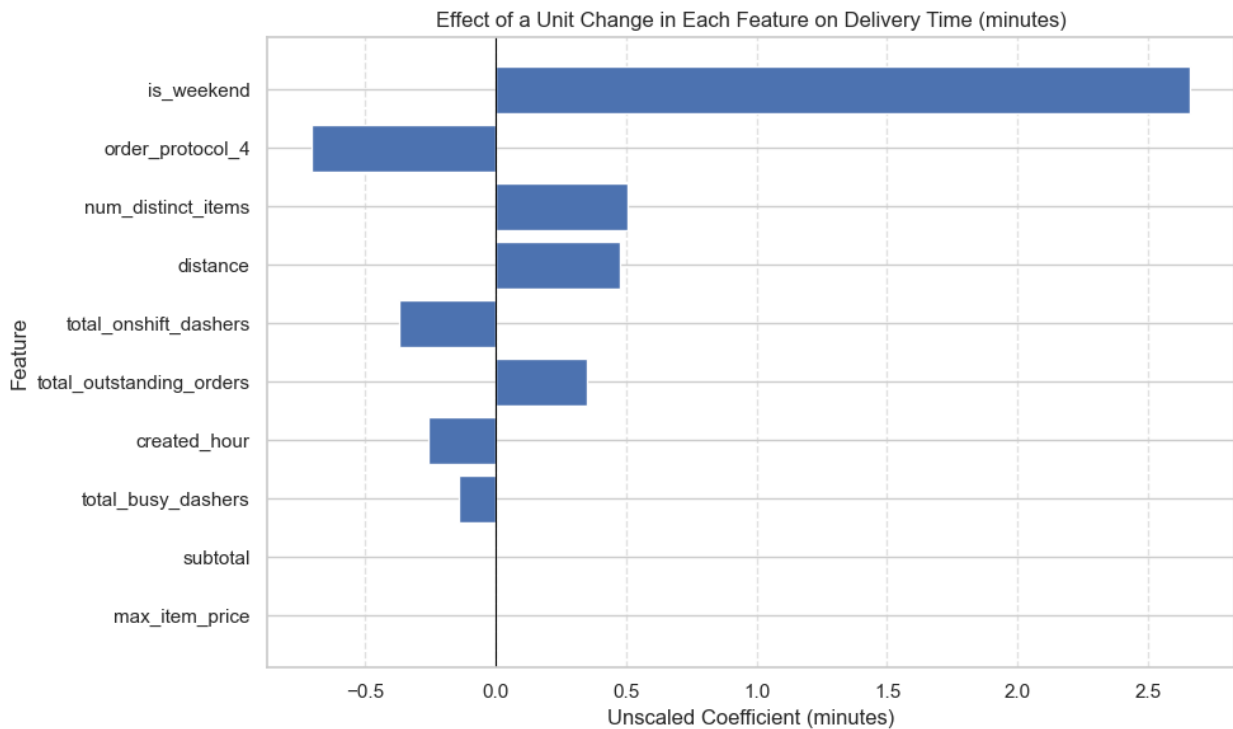
Additionally, we can analyse the effect of a unit change in a feature. In other words, because we have scaled the features, a unit change in the features will not translate directly to the model. Use scaled and unscaled coefficients to find how will a unit change in a feature affect the target.

```
# Analyze the effect of a unit change in a feature, say 'total_items'

coef_df["abs_coef"] = coef_df["Unscaled Coef"].abs()
df_sorted = coef_df.sort_values("abs_coef", ascending=True)

# Plot horizontal bar chart
plt.figure(figsize=(10, 6))
plt.barh(df_sorted["Feature"], df_sorted["Unscaled Coef"])
```

```
plt.axvline(0, color='black', linewidth=0.8)
plt.title("Effect of a Unit Change in Each Feature on Delivery Time (minutes)")
plt.xlabel("Unscaled Coefficient (minutes)")
plt.ylabel("Feature")
plt.grid(axis='x', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```



Note: The coefficients on the original scale might differ greatly in magnitude from the scaled coefficients, but they both describe the same relationships between variables.

Interpretation is key: Focus on the direction and magnitude of the coefficients on the original scale to understand the impact of each variable on the response variable in the original units.

Include conclusions in your report document.

## Subjective Questions [20 marks]

Answer the following questions only in the notebook. Include the visualisations/methodologies/insights/outcomes from all the above steps in your report.

### Subjective Questions based on Assignment

#### Question 1. [2 marks]

Are there any categorical variables in the data? From your analysis of the categorical variables from the dataset, what could you infer about their effect on the dependent variable?

**Answer:**

Yes, the dataset contains several categorical variables.

- weekday or weekend: weekends have longer delivery than weekdays
  - Market Id: some market location has higher average delivery time than others
  - store\_primary\_category: had minimal effect on the delivery time
- 

**Question 2.** [1 marks]

What does `test_size = 0.2` refer to during splitting the data into training and test sets?

**Answer:**

having `test_size 0.2` means splitting the data set into 80% and 20%; 80% of the data goes to training and 20% of the data goes to testing

---

**Question 3.** [1 marks]

Looking at the heatmap, which one has the highest correlation with the target variable?

**Answer:**

distance has the highest correlation with delivery time (`mints_taken`) with value of 0.46

---

**Question 4.** [2 marks]

What was your approach to detect the outliers? How did you address them?

**Answer:**

use IQR method on the outliers, i took 3 times of the values than normal becoz most of the potential outliers are not outliers

---

**Question 5.** [2 marks]

Based on the final model, which are the top 3 features significantly affecting the delivery time?

**Answer:**

`is_weekend`, `order_protocol_4` and `distance` are the three top features affecting the delivery time

---

## General Subjective Questions

### Question 6. [3 marks]

Explain the linear regression algorithm in detail

#### Answer:

- Linear regression is a supervised learning algorithm used to model the relationship between a dependent variable ( $y$ ) and one or more independent variables ( $X$ ).
  - Assume the relation between input and output is linear, constant variance, normality of errors
  - Each feature has coefficient which about the relation with the dependent variable.
- 

### Question 7. [2 marks]

Explain the difference between simple linear regression and multiple linear regression

#### Answer:

- simple linear regression has one independent variable, depends on single factor
  - multiple linear regression has more than one independent variable, depends on multiple factor
- 

### Question 8. [2 marks]

What is the role of the cost function in linear regression, and how is it minimized?

#### Answer:

- also known as loss function, quantifies the difference between predicted values and actual values.
  - lower the value the better the model.
  - The cost function guides the learning process to adjust coefficients to minimize prediction errors.
- 

### Question 9. [2 marks]

Explain the difference between overfitting and underfitting.

#### Answer:

- under fit: model is too simple to capture the data; high error on training and test set.
  - over fit: model is too complex, fits noise in the data, low training error and high testing error
-

**Question 10.** [3 marks]

How do residual plots help in diagnosing a linear regression model?

**Answer:**

- residual plot tells us about the difference in actual value and predicted value in graphical
- to check assumption: linear relation, constant variance, independence, Normality of errors.
- Random scatter ( no patter) is a good fit model.